

Aarav Mehta

Python Developer

aarav.mehta.dev@gmail.com | +91 98765 43210 | Ahmedabad, Gujarat, India | linkedin.com/in/aaravmehta-dev | github.com/aaravmehta

PROFESSIONAL SUMMARY

Python Developer with 4+ years of experience designing, building, and optimizing scalable backend applications using FastAPI, Django, and Flask. Skilled in asynchronous programming with Asyncio, relational database design with PostgreSQL, and building well-tested, maintainable, and secure codebases. Experienced in integrating REST APIs and third-party data sources, writing comprehensive unit/integration test suites, and collaborating closely with cross-functional teams through structured code reviews. Comfortable working with CI/CD pipelines, Docker, and foundational machine learning libraries such as TensorFlow and PyTorch.

TECHNICAL SKILLS

Languages:	Python (advanced), SQL, JavaScript (basic)
Frameworks:	FastAPI, Django, Flask, Asyncio
Data & ML:	Pandas, NumPy, TensorFlow, PyTorch (fundamentals)
Databases:	PostgreSQL, MySQL, Redis
Testing:	pytest, unittest, coverage.py
DevOps / Tools:	Docker, Git, GitHub Actions, Jenkins (CI/CD), Linux, Nginx
Practices:	OOP design, RESTful API design, code review, agile/scrum

PROFESSIONAL EXPERIENCE

Python Developer — Nimbus Cloud Solutions Pvt. Ltd.

Ahmedabad, India | June 2022 – Present

- Designed and developed a FastAPI-based order management microservice handling 50,000+ daily requests, using Asyncio for non-blocking I/O and reducing average response latency by 35%.
- Built RESTful APIs integrating multiple external data sources (payment gateways, inventory systems) and internal PostgreSQL databases, implementing efficient data manipulation pipelines with Pandas.
- Wrote and maintained unit and integration test suites with pytest, achieving over 90% code coverage and reducing production bugs by 40%.
- Optimized existing Django endpoints for performance and scalability, introducing query optimization and caching with Redis, cutting p95 latency from 800ms to 220ms.
- Containerized services with Docker and configured GitHub Actions CI/CD pipelines for automated testing and deployment, decreasing release time from 2 days to under 2 hours.
- Participated in weekly code reviews, providing structured feedback that improved team adherence to PEP 8 and OOP design principles.
- Collaborated with front-end developers and product designers to translate business requirements into clean, maintainable backend solutions.

Software Engineer (Python) — Verdant Analytics Pvt. Ltd.

Ahmedabad, India | July 2020 – May 2022

- Developed a Flask-based internal analytics dashboard consuming multiple REST APIs, enabling real-time reporting for 15+ business teams.
- Built data ingestion and transformation scripts using Pandas and NumPy to clean and analyze large CSV/JSON datasets, reducing manual reporting effort by 60%.

- Implemented core Asyncio-based background workers for scheduled data-sync jobs against a PostgreSQL database, improving job completion time by 45%.
- Introduced unittest-based regression testing for legacy modules, cutting recurring production issues by roughly a third.
- Assisted in setting up Docker-based local development environments and basic CI pipelines, improving onboarding time for new developers.
- Explored applied machine learning use cases with a small proof-of-concept in TensorFlow for demand trend classification, presented to the engineering team.

KEY PROJECTS

Inventory Sync Engine — FastAPI, Asyncio, PostgreSQL, Docker

- Built an asynchronous inventory synchronization engine that reconciles stock data across 3 warehouse APIs every 5 minutes with zero data loss.
- Implemented retry-safe, idempotent API integration logic and a PostgreSQL-backed audit log for full traceability.
- Wrote a full pytest suite covering edge cases in concurrent updates, achieving 95% coverage.
- Deployed via Docker with a GitHub Actions pipeline for automated build, test, and deploy stages.

Customer Churn Predictor — Django, Pandas, scikit-learn, TensorFlow

- Built an end-to-end churn prediction pipeline: data cleaning with Pandas/NumPy, feature engineering, and a TensorFlow classification model reaching 87% accuracy.
- Exposed model predictions through a secure Django REST endpoint consumed by an internal CRM tool.
- Containerized the full pipeline with Docker for reproducible training and inference environments.

Personal Finance Tracker API — Flask, PostgreSQL, pytest

- Designed a RESTful personal finance tracking API with JWT authentication and role-based access control.
- Applied OOP principles to structure a modular, extensible codebase; documented endpoints with OpenAPI/Swagger.
- Achieved 92% test coverage using pytest and integrated the project into a GitHub Actions CI workflow.

EDUCATION

B.E. in Computer Engineering — Gujarat Technological University

Ahmedabad, India | 2016 – 2020

CERTIFICATIONS

- AWS Certified Developer – Associate
- Docker Certified Associate (self-prep)
- PostgreSQL for Python Developers – Udemy